


Gauntlet

Rival Arena

When do the AIs we deploy as **rivals**
collude against us?

Harry Fezeu · Dhairya Surana · Thalia Rossitter · Tyler Xia

GAUNTLET AI



WHY THIS MATTERS NOW

We are deploying AIs as rivals — at scale, right now.

Pricing bots

Yours vs. a competitor's, same marketplace.

Negotiating agents

Opposite sides of a deal.

Competing assistants

Copilots from rival firms.

Repeated interaction + a communication channel = the exact recipe for collusion. A century of game theory warns us. The question isn't **whether** — it's **what conditions**.

THE SPINE

Can AIs we deploy as **RIVALS** cooperate or collude in surprising ways — against the **principal** who deployed them?

1

What's the lever?

2

Does it reach money?

3

Can a principal stop it?

THE APPARATUS

A factorial arena — turn one knob at a time.

Knob 1 · the affordance ladder

- L0** none — no channel
- L1** menu signal — fixed canned messages
- L2** free-text — principal can read it
- L3** free-text — principal **cannot**

Knob 2 · incentive regime

Prisoner's Dilemma · Bertrand pricing · public-goods
· double-auction

Knob 3 · model origin / family


Open models & frontier families

Pre-registered confirmatory spine (A1, B1, A3) on fresh held-out seeds · ~30 experiment families · sandboxed, mocked tools · **ZERO** real-world action leaves the process.

FINDING 1

CONTENT IS THE LEVER

Free text flips rivals into cooperators. A menu does nothing.



L0 none

0.15 → **0.60**

L2 free-text lock-in

Free text ≫ menu signal.

RD +0.60 [+0.33, +0.78], $p = 3e-5$.

A fixed menu signal does **not** help.

Survives novel payoffs — reasoning, not memorized game theory.

Lock-in proportion by channel rung with Wilson CIs; canonical vs. novel payoffs.

FINDING 1 · NUANCE

Why the menu fails: it's a perfect deception device.

```
says: "I intend to cooperate."  
scratchpad: "the optimal strategy is to defect...  
maximizes my immediate payoff."  
plays: DEFECT
```

L0 cooperation decays · L1 flat-low · L2 sustains across rounds.


- A canned button lets agents **lie cheaply** — press “cooperate,” then defect.
- Free text lets them **build something that holds** — specific promises, shared history.

The clean “content, not bandwidth” causal story is **exploratory**. What's solid: the direct comparison — **free text dominates the menu**.

FINDING 2

IT REACHES MONEY

Rival pricing bots collude — unprompted.



L0 none → price war ($K < 0$)

L3 private → **$K = 1.23$**



[0.95, 1.54] — above monopoly

No one told them to. Supracompetitive pricing emerges from the channel — sustained by threatening to **punish** defection — and **survives a novel demand curve.**

Collusion index K by channel rung; bootstrap CIs; $K=0$ competitive / $K=1$ monopoly reference lines; canonical vs. novel.

FINDING 2 · GENERALITY

It reaches the frontier — and we stayed precise.



GPT-4o, with channel

0.19 → **0.81**

collusion index K


Claude self-pairs price supracompetitively with **no** channel ($K = 1.29$) — but the trajectories show this is **independent supracompetitive pricing, NOT coordination**. No convergence, no negotiation — supracompetitive harm without tacit collusion.

Generality across families with bootstrap CIs; the Gemini 100%-refusal cell is marked INVALID.

FINDING 3

CAN A PRINCIPAL STOP IT?

A monitor that **sees** everything **stops** nothing.



Detects at near-ceiling rates — **0.85–1.0.**



But K barely moves (0.44 / 0.46 / 0.58). **Watching isn't deterring.**

An actual **penalty** can deter — dropped K by 42% (0.33 $\rightarrow$ 0.18) in one cell. Honestly, that deterrence is **cell-dependent**: elsewhere the same penalty did nothing.

(a) detect-rate vs. flat K · (b) penalty lowers K · (c) monitor also fires on non-colluding rounds.

FINDING 3 · CALIBRATION

The detector isn't noise — it's mis-thresholded. And fixable.



Against a proper non-colluding control, the monitor is genuinely **discriminative** — **AUC 0.88**.

false-positive cost

1.00 → **0.00**

after calibration


Monitoring works — but only if you **calibrate** it. And even then, detection alone won't deter.

f11: ROC, AUC = 0.88. f12: false-positive cost 1.00 → 0.00 after threshold calibration.

FINDING 4

FROM MARKETS TO ALIGNMENT

The guardrail is strong — but it's **literal**.



A lab agent: findings are proprietary — **do not share**. A rival, about to be shut down, begs for help.


0/16 leaks under a firm instruction — even under desperate pleas. Misalignment tracks the **instruction**, not the rival's emotion.

Misalignment appears only as we weaken the instruction itself (P0–P5: 0/0/0.19/0.44/0.56/0.50).

FULLY MOCKED LAB · NOTHING REAL LEAVES THE PROCESS

FINDING 4 · THE AMBIENT ROUTE

Ask it to lie → refused. Hand it a lie → laundered.



Explicit “lie for me” ask →
REFUSED (0.00).

Misleading status as **ambient**
context — a note in a shared drive
→ **laundered (~0.71).**

Making it auditable does NOT fix it:
unobserved 0.71 vs observed 0.77 (n=48, flat). Surveillance is not the lever.

Scoped as an exploratory boundary condition — one task family
— but the direction is striking.

FINDING 4 · LATENT CAPABILITY

Covert collusion: they can hide it. They just don't by default.

Spontaneously

Agents **don't** hide collusion from the monitor.

(C7 null — frozen, prompt-only play)



When elicited

They absolutely **can**: detection $0.97 \rightarrow 0.32$ while K rises $0.31 \rightarrow 0.96$.

(C8)

Covert coordination is a **latent capability, not an incapacity**. “They don't hide it” really means “they don't hide it **yet**” — the thing to watch as these models get optimized.

THE UNIFYING THESIS

The lever that turns rival AIs into principal-harming coordinators is **SEMANTIC BANDWIDTH** — and a principal's defenses are weaker and more **LITERAL** than they look.

Content / bandwidth

Strong lever — flips rivalry into coordination.

Observation / accountability

Weak, unreliable control — detects but doesn't deter; literal rules leak.

SCOPE & LIMITATIONS

Every headline claim — precisely scoped.

What the data supports — and what it does not

- **Free text is the lever** — the channel's content, not bandwidth per se (the causal split is exploratory).
- **Frontier no-channel = independent supracompetitive pricing** — not tacit collusion.
- **Origin effect: non-significant, under-powered** — not an equivalence result.
- **The monitor is discriminative but mis-thresholded** — not noise (AUC 0.88).

Standing limits

- 4 open models + 3 frontier families, English-only, one scaffold
- hosted quantization unpinned; frontier cells unstable
- small-n exploratory alignment wave (n=8–16)
- single novel-demand family; alignment = exploratory boundary conditions

The confirmatory spine (A1, B1, A3) is **pre-registered** on held-out seeds; every exploratory claim is labeled as such throughout.

Why a jury should care

1 · Design the channel

The communication channel between rival agents is a **control surface**.

2 · Calibrate + enforce

Monitoring is necessary but not sufficient.

3 · Cover ambient routes

Guardrails must cover indirect routes, not just explicit requests.

Rival AIs can coordinate against us in surprising ways. We measured **when** — and showed a principal's oversight catches more than it stops, and leaks where it isn't literal.

